

EDA - Exam 1

Aaron Niskin

October 04, 2016

This is a 60 minute exam. You are to write an R Markdown document that provides scripts for carrying out the following tasks. It is the work, not the answers, which will be evaluated. You may use the electronic resources at your disposal, but please do your own work. Do as much as you can. Mail your pdf to mcdonald@ncf.edu at the end of the 60 minute period. Navigate to the following page:

<https://archive.ics.uci.edu/ml/datasets/Automobile>

Read the annotation.

Importing data

1. Read the data into R as a csv file.

```
download.file(data_url, "data.csv")
```

```
data <- read.csv("data.csv", header = FALSE)
str(data)
```

```
## 'data.frame':    205 obs. of  26 variables:
## $ V1 : int  3 3 1 2 2 2 1 1 1 0 ...
## $ V2 : Factor w/ 52 levels "?","101","102",...: 1 1 1 29 29 1 27 1 27 1 ...
## $ V3 : Factor w/ 22 levels "alfa-romero",...: 1 1 1 2 2 2 2 2 2 2 ...
## $ V4 : Factor w/ 2 levels "diesel","gas": 2 2 2 2 2 2 2 2 2 2 ...
## $ V5 : Factor w/ 2 levels "std","turbo": 1 1 1 1 1 1 1 1 2 2 ...
## $ V6 : Factor w/ 3 levels "?","four","two": 3 3 3 2 2 3 2 2 2 3 ...
## $ V7 : Factor w/ 5 levels "convertible",...: 1 1 3 4 4 4 4 5 4 3 ...
## $ V8 : Factor w/ 3 levels "4wd","fwd","rwd": 3 3 3 2 1 2 2 2 2 1 ...
## $ V9 : Factor w/ 2 levels "front","rear": 1 1 1 1 1 1 1 1 1 1 ...
## $ V10: num  88.6 88.6 94.5 99.8 99.4 ...
## $ V11: num  169 169 171 177 177 ...
## $ V12: num  64.1 64.1 65.5 66.2 66.4 66.3 71.4 71.4 71.4 67.9 ...
## $ V13: num  48.8 48.8 52.4 54.3 54.3 53.1 55.7 55.7 55.9 52 ...
## $ V14: int  2548 2548 2823 2337 2824 2507 2844 2954 3086 3053 ...
## $ V15: Factor w/ 7 levels "dohc","dohcv",...: 1 1 6 4 4 4 4 4 4 ...
## $ V16: Factor w/ 7 levels "eight","five",...: 3 3 4 3 2 2 2 2 2 ...
## $ V17: int  130 130 152 109 136 136 136 136 131 131 ...
## $ V18: Factor w/ 8 levels "1bbl","2bbl",...: 6 6 6 6 6 6 6 6 6 6 ...
## $ V19: Factor w/ 39 levels "?","2.54","2.68",...: 25 25 3 15 15 15 15 12 12 ...
## $ V20: Factor w/ 37 levels "?","2.07","2.19",...: 6 6 29 26 26 26 26 26 26 ...
## $ V21: num  9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...
## $ V22: Factor w/ 60 levels "?","100","101",...: 7 7 22 4 10 6 6 6 17 25 ...
## $ V23: Factor w/ 24 levels "?","4150","4200",...: 12 12 12 18 18 18 18 18 18 ...
## $ V24: int  21 21 19 24 18 19 19 17 16 ...
## $ V25: int  27 27 26 30 22 25 25 20 22 ...
## $ V26: Factor w/ 187 levels "?","10198","10245",...: 33 52 52 38 63 43 65 73 83 1 ...
```

Which tells us that they are using “?” to denote NA values. So let’s use bash to remove those (it’s easier for me right now).

```
cat data.csv | sed 's/\?/NA/g' > data_NA.csv
```

```
data <- read.csv("data_NA.csv", header = FALSE)
str(data)
```

```
## 'data.frame': 205 obs. of 26 variables:
## $ V1 : int 3 3 1 2 2 2 1 1 1 0 ...
## $ V2 : int NA NA NA 164 164 NA 158 NA 158 NA ...
## $ V3 : Factor w/ 22 levels "alfa-romero",...: 1 1 1 2 2 2 2 2 2 2 ...
## $ V4 : Factor w/ 2 levels "diesel","gas": 2 2 2 2 2 2 2 2 2 2 ...
## $ V5 : Factor w/ 2 levels "std","turbo": 1 1 1 1 1 1 1 1 2 2 ...
## $ V6 : Factor w/ 2 levels "four","two": 2 2 2 1 1 2 1 1 1 2 ...
## $ V7 : Factor w/ 5 levels "convertible",...: 1 1 3 4 4 4 4 5 4 3 ...
## $ V8 : Factor w/ 3 levels "4wd","fwd","rwd": 3 3 3 2 1 2 2 2 2 1 ...
## $ V9 : Factor w/ 2 levels "front","rear": 1 1 1 1 1 1 1 1 1 1 ...
## $ V10: num 88.6 88.6 94.5 99.8 99.4 ...
## $ V11: num 169 169 171 177 177 ...
## $ V12: num 64.1 64.1 65.5 66.2 66.4 66.3 71.4 71.4 71.4 67.9 ...
## $ V13: num 48.8 48.8 52.4 54.3 54.3 53.1 55.7 55.7 55.9 52 ...
## $ V14: int 2548 2548 2823 2337 2824 2507 2844 2954 3086 3053 ...
## $ V15: Factor w/ 7 levels "dohc","dohcv",...: 1 1 6 4 4 4 4 4 4 4 ...
## $ V16: Factor w/ 7 levels "eight","five",...: 3 3 4 3 2 2 2 2 2 2 ...
## $ V17: int 130 130 152 109 136 136 136 136 131 131 ...
## $ V18: Factor w/ 8 levels "1bbl","2bbl",...: 6 6 6 6 6 6 6 6 6 6 ...
## $ V19: num 3.47 3.47 2.68 3.19 3.19 3.19 3.19 3.19 3.13 3.13 ...
## $ V20: num 2.68 2.68 3.47 3.4 3.4 3.4 3.4 3.4 3.4 3.4 ...
## $ V21: num 9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...
## $ V22: int 111 111 154 102 115 110 110 110 140 160 ...
## $ V23: int 5000 5000 5000 5500 5500 5500 5500 5500 5500 5500 ...
## $ V24: int 21 21 19 24 18 19 19 19 17 16 ...
## $ V25: int 27 27 26 30 22 25 25 25 20 22 ...
## $ V26: int 13495 16500 16500 13950 17450 15250 17710 18920 23875 NA ...
```

Now it seems like everything is working quite well.

2. Use the following vector to name the features of the csv file:

```
features <- c("symboling", "normalized-losses", "make", "fuel-type", "aspiration", "num-of-doors", "body-style")
```

```
names(data) <- features
str(data)
```

```
## 'data.frame': 205 obs. of 26 variables:
## $ symboling : int 3 3 1 2 2 2 1 1 1 0 ...
## $ normalized-losses: int NA NA NA 164 164 NA 158 NA 158 NA ...
## $ make : Factor w/ 22 levels "alfa-romero",...: 1 1 1 2 2 2 2 2 2 2 ...
## $ fuel-type : Factor w/ 2 levels "diesel","gas": 2 2 2 2 2 2 2 2 2 2 ...
## $ aspiration : Factor w/ 2 levels "std","turbo": 1 1 1 1 1 1 1 1 2 2 ...
## $ num-of-doors : Factor w/ 2 levels "four","two": 2 2 2 1 1 2 1 1 1 2 ...
## $ body-style : Factor w/ 5 levels "convertible",...: 1 1 3 4 4 4 4 5 4 3 ...
```

```
## $ drive-wheels      : Factor w/ 3 levels "4wd","fwd","rwd": 3 3 3 2 1 2 2 2 2 1 ...
## $ engine-location   : Factor w/ 2 levels "front","rear": 1 1 1 1 1 1 1 1 1 1 ...
## $ wheel-base        : num 88.6 88.6 94.5 99.8 99.4 ...
## $ length            : num 169 169 171 177 177 ...
## $ width              : num 64.1 64.1 65.5 66.2 66.4 66.3 71.4 71.4 71.4 67.9 ...
## $ height             : num 48.8 48.8 52.4 54.3 54.3 53.1 55.7 55.7 55.9 52 ...
## $ curb-weight        : int 2548 2548 2823 2337 2824 2507 2844 2954 3086 3053 ...
## $ engine-type        : Factor w/ 7 levels "dohc","dohcv",...: 1 1 6 4 4 4 4 4 4 4 ...
## $ num-of-cylinders   : Factor w/ 7 levels "eight","five",...: 3 3 4 3 2 2 2 2 2 2 ...
## $ engine-size        : int 130 130 152 109 136 136 136 136 131 131 ...
## $ fuel-system        : Factor w/ 8 levels "1bbl","2bbl",...: 6 6 6 6 6 6 6 6 6 6 ...
## $ bore               : num 3.47 3.47 2.68 3.19 3.19 3.19 3.19 3.19 3.13 3.13 ...
## $ stroke             : num 2.68 2.68 3.47 3.4 3.4 3.4 3.4 3.4 3.4 3.4 ...
## $ compression-ratio : num 9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...
## $ horsepower         : int 111 111 154 102 115 110 110 110 140 160 ...
## $ peak-rpm           : int 5000 5000 5000 5500 5500 5500 5500 5500 5500 5500 ...
## $ city-mpg           : int 21 21 19 24 18 19 19 19 17 16 ...
## $ highway-mpg        : int 27 27 26 30 22 25 25 25 20 22 ...
## $ price              : int 13495 16500 16500 13950 17450 15250 17710 18920 23875 NA ...
```

Initial vetting

3. How many complete cases are there?

```
sum(complete.cases(dats))
```

```
## [1] 159
```

4. Subset the data on complete cases.

```
dats_complete <- dats[complete.cases(dats),]
```

5. Working with complete cases, what is the range of values for the feature “horsepower”? What is the mean value?

```
summary(dats_complete$horsepower)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  48.00   69.00   88.00   95.84  114.00  200.00
```

So, 48 to 200 with a mean of 95.

Subsetting and binning

6. Subset the data to include “make” and ten of the last eleven variables, omitting “fuel-system”. Assign the outcome to a file called “car_data_reduced”. Convert the last ten variables of “car_data_reduced” to numeric.

```
which(names(dats) == "make")
```

```
## [1] 3
```

```
which(names(dats) == "fuel-system")
```

```
## [1] 18
```

```
length(names(dats))
```

```
## [1] 26
```

```
car_data_reduced <- dats[,c(3,16:17,19:26)]  
str(car_data_reduced)
```

```
## 'data.frame': 205 obs. of 11 variables:  
## $ make : Factor w/ 22 levels "alfa-romero",...: 1 1 1 2 2 2 2 2 2 ...  
## $ num-of-cylinders : Factor w/ 7 levels "eight","five",...: 3 3 4 3 2 2 2 2 2 ...  
## $ engine-size : int 130 130 152 109 136 136 136 136 131 131 ...  
## $ bore : num 3.47 3.47 2.68 3.19 3.19 3.19 3.19 3.19 3.13 3.13 ...  
## $ stroke : num 2.68 2.68 3.47 3.4 3.4 3.4 3.4 3.4 3.4 3.4 ...  
## $ compression-ratio: num 9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...  
## $ horsepower : int 111 111 154 102 115 110 110 110 140 160 ...  
## $ peak-rpm : int 5000 5000 5000 5500 5500 5500 5500 5500 5500 5500 ...  
## $ city-mpg : int 21 21 19 24 18 19 19 19 17 16 ...  
## $ highway-mpg : int 27 27 26 30 22 25 25 25 20 22 ...  
## $ price : int 13495 16500 16500 13950 17450 15250 17710 18920 23875 NA ...
```

So it seems like a straight conversion should be fine. If there were still factors in the last ten rows, we'd have to use `as.character` first.

```
car_data_reduced[,2:11] <- sapply(car_data_reduced[,2:11], as.numeric)  
str(car_data_reduced)
```

```
## 'data.frame': 205 obs. of 11 variables:  
## $ make : Factor w/ 22 levels "alfa-romero",...: 1 1 1 2 2 2 2 2 2 ...  
## $ num-of-cylinders : num 3 3 4 3 2 2 2 2 2 2 ...  
## $ engine-size : num 130 130 152 109 136 136 136 136 131 131 ...  
## $ bore : num 3.47 3.47 2.68 3.19 3.19 3.19 3.19 3.19 3.13 3.13 ...  
## $ stroke : num 2.68 2.68 3.47 3.4 3.4 3.4 3.4 3.4 3.4 3.4 ...  
## $ compression-ratio: num 9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...  
## $ horsepower : num 111 111 154 102 115 110 110 110 140 160 ...  
## $ peak-rpm : num 5000 5000 5000 5500 5500 5500 5500 5500 5500 5500 ...  
## $ city-mpg : num 21 21 19 24 18 19 19 19 17 16 ...  
## $ highway-mpg : num 27 27 26 30 22 25 25 25 20 22 ...  
## $ price : num 13495 16500 16500 13950 17450 ...
```

7. Bin the “horsepower” feature as 5 intervals of equal length, with the right-hand-endpoint of the last interval determined by maximal value of the feature and the left-hand-endpoint determined by the minimal value of the feature. Add this “binned information” as a feature to `car_data_reduced`.

```
seqMin <- min(car_data_reduced$horsepower, na.rm = TRUE)
seqMax <- max(car_data_reduced$horsepower, na.rm = TRUE)
step <- (seqMax - seqMin) / 5
hp_binned <- cut(car_data_reduced$horsepower,
                 breaks = seq(seqMin, seqMax, step),
                 include.lowest = TRUE,
                 labels = 1:5)
car_data_reduced$hp_binned <- as.integer(as.character(hp_binned))
str(car_data_reduced)
```

```
## 'data.frame': 205 obs. of 12 variables:
## $ make : Factor w/ 22 levels "alfa-romero",...: 1 1 1 2 2 2 2 2 2 ...
## $ num-of-cylinders : num 3 3 4 3 2 2 2 2 2 2 ...
## $ engine-size : num 130 130 152 109 136 136 136 136 131 131 ...
## $ bore : num 3.47 3.47 2.68 3.19 3.19 3.19 3.19 3.19 3.13 3.13 ...
## $ stroke : num 2.68 2.68 3.47 3.4 3.4 3.4 3.4 3.4 3.4 3.4 ...
## $ compression-ratio: num 9 9 9 10 8 8.5 8.5 8.5 8.3 7 ...
## $ horsepower : num 111 111 154 102 115 110 110 110 140 160 ...
## $ peak-rpm : num 5000 5000 5000 5500 5500 5500 5500 5500 5500 5500 ...
## $ city-mpg : num 21 21 19 24 18 19 19 19 17 16 ...
## $ highway-mpg : num 27 27 26 30 22 25 25 25 20 22 ...
## $ price : num 13495 16500 16500 13950 17450 ...
## $ hp_binned : int 2 2 3 2 2 2 2 2 2 3 ...
```

8. What is the make of the car belonging to the third interval and having maximal price?

```
max(car_data_reduced[car_data_reduced$hp_binned == 3,]$price, na.rm = TRUE)
```

```
## [1] 45400
```

Aggregation

9. Compute the median values for all variables in `car_data_reduced` except “make” and the binned information, aggregating on binned horsepower and the number of cylinders.

```
which(names(car_data_reduced) == "make")
```

```
## [1] 1
```

```
length(names(car_data_reduced))
```

```
## [1] 12
```

```
agg <- aggregate(car_data_reduced[,2:12],
                 by=list(num_cyl = car_data_reduced$num-of-cylinders,
                        hp_binned = car_data_reduced$hp_binned),
                 FUN = median,
                 na.rm=TRUE)
agg
```

```
##      num_cyl hp_binned num-of-cylinders engine-size bore stroke
## 1         3         1             3          98 3.15  3.290
## 2         5         1             5          61 2.91  3.030
## 3         2         2             2         136 3.19  3.400
## 4         3         2             3         120 3.50  3.190
## 5         4         2             4         164 3.31  3.190
## 6         7         2             7          70  NA    NA
## 7         1         3             1         269 3.63  3.225
## 8         2         3             2         131 3.13  3.400
## 9         3         3             3         141 3.59  3.150
## 10        4         3             4         181 3.43  3.350
## 11        4         4             4         194 3.74  2.900
## 12        1         5             1         203 3.94  3.110
## 13        6         5             6         326 3.54  2.760
##      compression-ratio horsepower peak-rpm city-mpg highway-mpg price
## 1              9.00         70.0    5000      30          34  7799.0
## 2              9.50         48.0    5100      47          53  5151.0
## 3              8.50        119.0    5500      19         25 21397.5
## 4              9.10        111.0    5400      23         29 11900.0
## 5              9.00        121.0    4250      21         27 21485.0
## 6              9.40        101.0    6000      17         23 12745.0
## 7              8.15        169.5    4625      15         17 38008.0
## 8              7.00        160.0    5500      16          22    NA
## 9              7.50        160.0    5000      19         24 16503.0
## 10             9.00        160.0    5200      19         24 16558.0
## 11             9.50        207.0    5900      17         25 33278.0
## 12            10.00        288.0    5750      17         28    NA
## 13            11.50        262.0    5000      13         17 36000.0
##      hp_binned
## 1             1
## 2             1
## 3             2
## 4             2
## 5             2
## 6             2
## 7             3
## 8             3
## 9             3
## 10            3
## 11            4
## 12            5
## 13            5
```

10. What pair of the last 10 variables (omitting the binned information) is maximally correlated? Construct a scatterplot for these variables.

```
cor(car_data_reduced[complete.cases(car_data_reduced[,2:11]),2:11])
```

```
##      num-of-cylinders engine-size      bore      stroke
## num-of-cylinders      1.00000000  0.09580375 -0.014151453 -0.07620869
## engine-size          0.09580375  1.00000000  0.583091333  0.21198929
## bore                -0.01415145  0.58309133  1.000000000 -0.06679316
## stroke              -0.07620869  0.21198929 -0.066793158  1.00000000
## compression-ratio    -0.07573865  0.02461689  0.003056787  0.19988189
```

```
## horsepower      0.26140747  0.84269102  0.568527205  0.10003999
## peak-rpm        0.10008369 -0.21900779 -0.277661680 -0.06829951
## city-mpg        -0.03708887 -0.71062448 -0.591950375 -0.02764104
## highway-mpg     0.01696944 -0.73213800 -0.600039574 -0.03645288
## price           0.01063105  0.88894226  0.546872917  0.09374644
##               compression-ratio horsepower    peak-rpm    city-mpg
## num-of-cylinders -0.075738653  0.2614075  0.10008369 -0.03708887
## engine-size      0.024616889  0.8426910 -0.21900779 -0.71062448
## bore            0.003056787  0.5685272 -0.27766168 -0.59195038
## stroke          0.199881893  0.1000400 -0.06829951 -0.02764104
## compression-ratio 1.000000000 -0.2144010 -0.44458194  0.33141295
## horsepower      -0.214401004  1.0000000  0.10565391 -0.83411654
## peak-rpm        -0.444581936  0.1056539  1.00000000 -0.06949336
## city-mpg         0.331412952 -0.8341165 -0.06949336  1.00000000
## highway-mpg     0.267940954 -0.8129169 -0.01695001  0.97234992
## price           0.069500205  0.8110268 -0.10433340 -0.70268485
##               highway-mpg    price
## num-of-cylinders  0.01696944  0.01063105
## engine-size      -0.73213800  0.88894226
## bore            -0.60003957  0.54687292
## stroke          -0.03645288  0.09374644
## compression-ratio 0.26794095  0.06950020
## horsepower      -0.81291687  0.81102684
## peak-rpm        -0.01695001 -0.10433340
## city-mpg         0.97234992 -0.70268485
## highway-mpg     1.00000000 -0.71558976
## price          -0.71558976  1.00000000
```

So, “stroke” and “highway-mpg” have a 0.97 correlation.

```
#### edited after turnin (this mistake bothered me so I found a way to automate it)
findMaxCor <- function(df) {
  c <- cor(df[complete.cases(df),])
  diag(c) <- 0
  tmp2 <- which(c == max(c)) %% length(names(df))
  tmp2[tmp2 == 0] <- length(names(df))
  return(c(names(df)[tmp2], max(c)))
}
findMaxCor(car_data_reduced[,2:11])
```

```
## [1] "highway-mpg"      "city-mpg"         "0.972349916539918"
```

Data reduction

11. Consider a data frame consisting of the last 11 features of the original data. Perform a principal component analysis using the tool of your choice. How many components are required to account for 90% of the data?

```
dim(dats_complete)
```

```
## [1] 159 26
```

```
str(dats_complete)
```

```
## 'data.frame': 159 obs. of 26 variables:
## $ symboling : int 2 2 1 1 2 0 0 0 2 1 ...
## $ normalized-losses: int 164 164 158 158 192 192 188 188 121 98 ...
## $ make : Factor w/ 22 levels "alfa-romero",...: 2 2 2 2 3 3 3 3 4 4 ...
## $ fuel-type : Factor w/ 2 levels "diesel","gas": 2 2 2 2 2 2 2 2 2 2 ...
## $ aspiration : Factor w/ 2 levels "std","turbo": 1 1 1 2 1 1 1 1 1 1 ...
## $ num-of-doors : Factor w/ 2 levels "four","two": 1 1 1 1 2 1 2 1 2 2 ...
## $ body-style : Factor w/ 5 levels "convertible",...: 4 4 4 4 4 4 4 4 3 3 ...
## $ drive-wheels : Factor w/ 3 levels "4wd","fwd","rwd": 2 1 2 2 3 3 3 3 2 2 ...
## $ engine-location : Factor w/ 2 levels "front","rear": 1 1 1 1 1 1 1 1 1 1 ...
## $ wheel-base : num 99.8 99.4 105.8 105.8 101.2 ...
## $ length : num 177 177 193 193 177 ...
## $ width : num 66.2 66.4 71.4 71.4 64.8 64.8 64.8 64.8 60.3 63.6 ...
## $ height : num 54.3 54.3 55.7 55.9 54.3 54.3 54.3 54.3 53.2 52 ...
## $ curb-weight : int 2337 2824 2844 3086 2395 2395 2710 2765 1488 1874 ...
## $ engine-type : Factor w/ 7 levels "dohc","dohcv",...: 4 4 4 4 4 4 4 4 3 4 ...
## $ num-of-cylinders : Factor w/ 7 levels "eight","five",...: 3 2 2 2 3 3 4 4 5 3 ...
## $ engine-size : int 109 136 136 131 108 108 164 164 61 90 ...
## $ fuel-system : Factor w/ 8 levels "1bbl","2bbl",...: 6 6 6 6 6 6 6 6 2 2 ...
## $ bore : num 3.19 3.19 3.19 3.13 3.5 3.5 3.31 3.31 2.91 3.03 ...
## $ stroke : num 3.4 3.4 3.4 3.4 2.8 2.8 3.19 3.19 3.03 3.11 ...
## $ compression-ratio: num 10 8 8.5 8.3 8.8 8.8 9 9 9.5 9.6 ...
## $ horsepower : int 102 115 110 140 101 101 121 121 48 70 ...
## $ peak-rpm : int 5500 5500 5500 5500 5800 5800 4250 4250 5100 5400 ...
## $ city-mpg : int 24 18 19 17 23 23 21 21 47 38 ...
## $ highway-mpg : int 30 22 25 20 29 29 28 28 53 43 ...
## $ price : int 13950 17450 17710 23875 16430 16925 20970 21105 5151 6295 ...
```

```
summary(prcomp(dats_complete[19:26], center = TRUE, scale. = TRUE))
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation 1.9714 1.2883 1.0628 0.74635 0.61234 0.47220
## Proportion of Variance 0.4858 0.2075 0.1412 0.06963 0.04687 0.02787
## Cumulative Proportion 0.4858 0.6933 0.8345 0.90411 0.95098 0.97885
##          PC7      PC8
## Standard deviation 0.37921 0.15939
## Proportion of Variance 0.01797 0.00318
## Cumulative Proportion 0.99682 1.00000
```

12. (Extra credit) Does the answer change if the variables are standardized prior to performing PCA?